

ISHAN DURUGKAR

@ ishand@cs.utexas.edu

+1 (919)987-5230

Austin, TX

<https://idurugkar.github.io/>

EDUCATION

PH.D. COMPUTER SCIENCE (2017+)

THE UNIVERSITY OF TEXAS AT AUSTIN, USA

Research areas: Reinforcement Learning, Multi-agent Reinforcement Learning, Robotics, Imitation Learning, Deep Learning

Affiliations: Learning Agents Research Group (LARG)

M.S. COMPUTER SCIENCE (2014 – 2017)

UNIVERSITY OF MASSACHUSETTS AMHERST

Research Areas: Deep Learning, Reinforcement Learning, Generative Adversarial Networks

GPA: 3.97

B.ENG. COMPUTER ENGINEERING (2007 – 2013)

MAHARASHTRA INSTITUTE OF TECHNOLOGY, INDIA

Research and Extracurricular interests: Robotics, Swarm techniques, non-convex optimization

BIO & RESEARCH STATEMENT

I am a PhD Student in the Department of Computer Science at the University of Texas at Austin. I am part of the Learning Agents Research Group (LARG) and advised by Prof. Peter Stone. My research focuses on the sub-field of Machine Learning called Reinforcement Learning. My doctoral thesis will focus on improving RL algorithms using the estimation and control of the visitation distributions induced by reinforcement learning agents, such as the distribution of states they visit, or the distribution of transitions they experience.

Additionally, I have worked on a variety of projects that look at different aspects of the RL problem. Some of these directions that I have worked on are optimization perspectives of TD learning, considering agent preferences in multi-agent RL, using adversarial techniques for transferring policies from simulation to real robots, and some applications of RL (for example, learning a policy to navigate knowledge bases).

Apart from these research interests, I also coordinate the weekly RL Reading Group with participation across the University. I also participate in the RoboCup Standard Platform League as part of the UT Austin Villa team. Here I work with the Nao robots, and focus on their vision system for the competition, along with using it as a testbed for research on sim-to-real transfer.

In Summer 2021, I interned with DeepMind, where I worked with Volodymyr Mnih on unsupervised RL. I have also interned with Microsoft Research in Redmond, WA, in the Summer of 2018, where I worked with Adith Swaminathan and Matthew Hausknecht on figuring out how to help developers train RL agents.

RESEARCH WORK EXPERIENCE

Research Assistant

THE UNIVERSITY OF TEXAS AT AUSTIN

2017 +

Assisted and authored research with UT Austin faculty and other graduate students on various projects. Primary research focus has been on Reinforcement Learning with different projects focusing on different aspects of this problem, including optimization, human feedback, sim-to-real transfer, and multi-agent RL.

Research Scientist Intern

DEEPMIND

Summer 2020

Research and development of deep reinforcement learning algorithms. Focus on unsupervised RL and the learning of skillful policies in the absence of specific learning signals from the environment. Evaluation on the Atari benchmark.

Research Intern

MICROSOFT RESEARCH, REDMOND

Summer 2018

Researched and implemented algorithms to improve ease of developing RL agents. Particularly, focused on incorporating multiple different feedback channels into learning and agent policy. Concretely tested the idea by leveraging demonstrations to speed up the A2C algorithm.

Research Assistant

UNIVERSITY OF MASSACHUSETTS AT AMHERST

2015–2017

Assisted and authored research with faculty and other graduate students on various research topics including: (1) Generative Adversarial Networks (GANs), (2) learning temporal abstractions for actions with Deep Q Networks, (3) manifold learning on graphs, and (4) learning policies to navigate large knowledge bases.

Research Member

MAHARASHTRA INSTITUTE OF TECHNOLOGY

2013

Participated in research on non-convex optimization using swarm techniques, leading to a novel Cohort Intelligence algorithm. Presented the research at IEEE Systems, Man and Cybernetics conference in 2013 at Manchester, UK.

OTHER WORK EXPERIENCE

Software Engineering Intern

AMAZON COMIXOLOGY

Summer 2015

Implemented certain components of the Comixology continuous integration pipeline. Developed a build system based on Ruby, PHP and ANT. Enabled developers to run Unit, Integration and Functional Tests in a single setting. Complete development of this system and integration into the development process.

Data Scientist – Research Assistant

UMASS AMHERST ALUMNI ASSOCIATION

2015 – 2016

Used data aggregation, cleaning and analysis along with machine learning techniques to identify and better engage alumni with on campus events and fund raising.

Software Engineer

MICROSOFT INDIA

2013 – 2014

As part of the Code Search team in Visual Studio Online, implemented a prototype of code search based on Elasticsearch, which was then approved for further development. Participated in the effort to scale the prototype to integrate with Visual Studio Online.

COMPUTING SKILLS

Python (proficient); C++ (intermediate); Experience with various NN training libraries (Tensorflow, PyTorch, Theano); HPC clusters; Distributed NN Training.

RELEVANT COURSEWORK

University of Texas at Austin: Reinforcement Learning, Advanced Robotics, Optimization.

University Of Massachusetts Amherst: Advanced Machine Learning, Deep Learning, Probabilistic Graphical Models, Optimization, Advanced Algorithms.

ROBOTICS COMPETITIONS

RoboCup

UT AUSTIN VILLA

2019 - 2020

Quarter-finalists in Standard Platform League

ABU Asia-Pacific Robot Contest

MAHARASHTRA INSTITUTE OF TECHNOLOGY

2010 – 2012

TEACHING

Teaching Assistant

THE UNIVERSITY OF TEXAS AT AUSTIN

Fall 2019

Reinforcement Learning: Theory and Practice

Teaching Assistant

UMASS AMHERST

2015 – 2017

- Applied Machine Learning
- Distributed and Operating Systems

REVIEWING EXPERIENCE

Academic Conferences:

- AAAI Conference on Artificial Intelligence 2021, 2022
- International Conference on Learning Representations (ICLR) 2021, 2022
- Conference on Neural Information Processing Systems (NeurIPS) 2020, 2021
- International Conference on Machine Learning (ICML) 2021, 2022
- International Conference on
- International Joint Conferences on Artificial Intelligence (IJCAI) 2020, 2021
- International Conference on Artificial Intelligence and Statistics (AISTATS) 2022

Journals:

- Journal of Artificial Intelligence Research (JAIR), 2022
- IEEE Transactions on Neural Networks and Learning Systems (TNNLS) 2020, 2021

Workshops:

- Lifelong Machine Learning (ICML 2020)
- Multi-Task and Lifelong Reinforcement Learning (ICML 2019)
- Neverending learning workshop (ICLR 2021)

ORGANIZATIONAL

RL Reading Group

UT AUSTIN

Coordinator

Spring 2019 -

Indian Student Association

UNIVERSITY OF MASSACHUSETTS, AMHERST

Fall 2015 – Fall 2016

Vice-President of the Indian Student Association at the University of Massachusetts, Amherst.

PUBLICATIONS

Conference Proceedings

- Durugkar, Ishan, Mauricio Tec, et al. (2021). “Adversarial Intrinsic Motivation for Reinforcement Learning”. In: *Advances in Neural Information Processing Systems*.
- Desai, Siddharth et al. (2020). “An Imitation from Observation Approach to Transfer Learning with Dynamics Mismatch”. In: *Advances in Neural Information Processing Systems*.
- Durugkar, Ishan, Elad Liebman, and Peter Stone (2020). “Balancing Individual Preferences and Shared Objectives in Multiagent Reinforcement Learning”. In: *Proceedings of the 29th International Joint Conference on Artificial Intelligence*, pp. 2505–2511.
- Pavse, Brahma et al. (2020). “Reducing Sampling Error in Batch Temporal Difference Learning”. In: *Proceedings of 37th International Conference on Machine Learning, Vienna, Austria*.
- Das, Rajarshi et al. (2018). “Go for a Walk and Arrive at the Answer: Reasoning Over Paths in Knowledge Bases using Reinforcement Learning”. In: *International Conference on Learning Representations*.
- Durugkar, Ishan, Ian Gemp, and Sridhar Mahadevan (2017). “Generative multi-adversarial networks”. In: *International Conference on Learning Representations*.
- Thomas, Philip S et al. (2017). “Predictive off-policy policy evaluation for nonstationary decision problems, with applications to digital marketing”. In: *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence*, pp. 4740–4745.
- Kulkarni, Anand J, Ishan P Durugkar, and Mrinal Kumar (2013). “Cohort intelligence: a self supervised learning behavior”. In: *2013 IEEE international conference on systems, man, and cybernetics*. IEEE, pp. 1396–1400.

Preprints

- Durugkar, Ishan, Steven Hansen, et al. (2021). *Wasserstein Distance Maximizing Intrinsic Control*. NeurIPS Deep Reinforcement Learning Workshop.
- Durugkar, Ishan, Matthew Hausknecht, et al. (2019). *Multi-Preference Actor Critic*. arXiv preprint arXiv:1904.03295.
- Durugkar, Ishan and Peter Stone (2018). *TD learning with constrained gradients*. NeurIPS Deep Reinforcement Learning Symposium.
- Durugkar, Ishan P et al. (2016). *Deep reinforcement learning with macro-actions*. arXiv preprint arXiv:1606.04615.